

Figure 2. Prompt Contract - 6-Field Schema for Every LLM Call

Each agent invocation must produce a payload that satisfies all six fields, enabling pre-generation policy enforcement and post-generation review.

task What the agent is asked to do. Verb-first, single-sentence intent. e.g. 'Generate a tactical comment for possession poss_p1_e20.'	source_scope Whitelisted evidence sources. PBP slice, RAG section IDs, video frame range. Anything outside this scope is forbidden.	evidence_requirements Per-assertion evidence binding. Every claim must cite an observation_id or fact_id. Unsupported claims are rejected.
forbidden_behaviors Hard-block list. e.g.: - inventing player names - injecting opinions on referees - forcing tactical labels onto non-tactical events	output_contract Strict output shape. JSON schema with required keys and value types. Parse failure triggers regeneration.	review_gate Downstream agents that may block or amend this output. Typically Fact Checker, then Risk Guard. Each gate has veto rights.

All six fields are validated by a schema gate before the LLM is invoked; fields are also re-checked against the generated output by Fact Checker and Risk Guard.